

General Exam Guidelines

- **Deadline:** 26 May 2026 (check on LearnIT)
- **Groups:** 2 - 3 people
- **What to submit:**
 1. Max 10 pages report (find template on LearnIT, under the Exam section).
References are not included as a part of the 10 page requirement.
 2. Jupyter notebook (well-curated and that runs without errors)
 3. Statement of contribution (find file on LearnIT, under the Exam section)
- **Mandatory Assignment(s):** Submit your mandatory assignment(s) 1 or 2 if you haven't submitted them during the course or they were not approved. Submit them individually (even if you have worked in groups).

Styling Requirements for the Report

We setup the LaTeX and Word templates according to these requirements:

1. **Citation Style:** [Vancouver Style](#) bibliography. The LaTeX template is already set up with the Vancouver Style. If you work in Word, you have to make sure you use the correct citation style: you can extract Vancouver Style citations using Google Scholar ([here is the guide](#)).
2. **Font:** Times New Roman. 12pt font size for body, 1.5pt line spacing, and 2cm margins (the templates already have this preset).

General Guidelines for the Report

Depending on your chosen project, not all questions are relevant/can be addressed. This is not necessarily an exhaustive list of questions; an important question relevant for your project might not be listed. Consider the feedback you got about your project (i.e. Mandatory Assignment 3) in this regard. **Please note, everybody needs to address the ethics/philosophy part (point 7 below).**

1. Description of the chosen problem.
 - Do you have access to the training data?
 - Do you have access to the algorithm?
 - Do you have access to the prediction?

-
- What is the stated purpose of the algorithm? Does it do so?
 - Why is this a problem, and who (e.g. which groups) does it affect?
 2. Dataset analysis: What are the characteristics of the dataset? Are there biases? What do you expect, given this analysis? Are there biases in the training data? Are there biases in the predictions?
 3. Is this algorithm fair? Which notion(s) of fairness did you use and why?
 - Motivate your choice(s) of fairness metric(s)
 - Which context is the algorithm used in?
 4. Can you explain the results of the algorithm? How can you explain them?
 - It is a blackbox, or whitebox algorithm. Which methods can be used to explain predictions?
 5. Can you debias the model and/or the data? If it's impossible, why is this impossible?
 6. How does fairness improve after debiasing? Does it even improve? Does your fairness approach work?
 7. Critical reflections. Were you able to improve the fairness of the system? Should such a system even be used? Reflect critically on your results, and, if relevant, provide recommendations to what e.g. the developers of the algorithm should do, what should the civil society do, etc.
 8. Ethics and philosophy (considerations on these two aspects):
 - Different ethical approaches and what could be said about dataset biases from the point of view of the theories that Pawel described during the lectures.
 - The switch from the paradigm of Good Old Fashioned Intelligence to modern AI, especially the arguments as to why datasets are used to solve specific problems. Were these problems unsolvable using traditional, algorithmic AI, or is it just a pragmatic choice (for example, these solutions are faster)?